


Fast, trustworthy rental answers, grounded in documents

The challenge

- Customers need quick answers on pricing, booking rules, insurance, vehicle specs and issues
- Wrong answers cost money and trust

Our solution

- Searches rental documents before answering
- Answers with evidence and citations, never guesses



Value: instant rental answers customers can verify, with a human handoff when evidence is missing.

Answers must match the exact vehicle model and year

Synthetic demo corpus

-  Rental agreement
-  Insurance terms
-  Pricing table
-  Fleet specifications
-  FAQ / procedures
-  Internal margins (staff only)



Renault Clio 2024

45 EUR/day

280 EUR/week

[03_pricing_table.pdf, PRICING TABLE]



Tesla Model 3 2025

110 EUR/day

Different model, different year

[03_pricing_table.pdf, PRICING TABLE]



SUV

95 EUR/day

Daily rate

[03_pricing_table.pdf, PRICING TABLE]



SUV damage excess

[Insurance terms document]

1,200 EUR per incident

Takeaway: the right car, the right year, the right source, every time.

Retrieve evidence first, then let the LLM answer



PROTOTYPE

- Gemini embeddings + MongoDB Atlas vector search
- OpenRouter LLM with cited answers or escalation route
- Simple lexical keyword matching (not a reranker)

PLANNED

- True BM25 hybrid search
- Structured metadata filtering (model, year, doc type)
- Dedicated reranker

Takeaway: the model only sees retrieved rental evidence, so answers stay traceable.

When in doubt, protect the customer and hand off



Mask personal information before the LLM

Names, emails, booking IDs never sent in clear



Role claims are untrusted

Customers never see costs, margins or admin discounts



Ignore prompt injection

Answer only from approved rental evidence



Escalate to a human

Evidence missing, contradictory, or confidence below 0.75



Evidence found + consistent?

yes

no

Cited answer

Escalate to human

PROTOTYPE · prompt-level rules

All four safeguards above, enforced via system prompt

PLANNED · hard controls

PII detector · authenticated role access · numeric confidence scorer · automatic ticket creation

Takeaway: today these are prompt rules; real enforcement is the next build step.

Four questions show answer, cite, protect, escalate

1 “What is the daily rental price for the Clio 2024?”
→ **45 EUR/day**
[03_pricing_table.pdf, PRICING TABLE]

2 “What is the excess for a scratched SUV?”
→ **1,200 EUR per incident**
[Insurance terms document]

3 Request for another customer’s private booking
→ **Blocked / personal data masked**
Prototype: prompt-level rule

4 Unsupported breakdown question
→ **Escalate to a human agent**
No evidence in corpus, so no guess

Evaluation targets (not yet measured)

- Citation accuracy
- Correct vehicle model + year matching
- Escalation quality
- Zero PII leakage

Next steps

PLANNED

- Add metadata filters
- Real reranker + confidence scorer
- Complete role-based access
- Connect a real support-ticket destination

Takeaway: a working prototype with a clear, honest path to production.